

# UNIX System Programming

## Signals & Resource Limits

### Module 3

# Agenda

---

**1**

**Signals**

**2**

**Resource Limits**

# Objectives

---

At the end of this module you will be able to:

- understand the purpose of signals
- handle signals in applications
- understand POSIX signal handling
- know resource limits defined by UNIX
- alter resource limits
- test the impact of resource limits on processes



# Signals



# Purpose of Signals

---

- Signal is a mechanism to indicate occurrence of an event to a process
- Signal is a software interrupt or notification.
- Signals help to manage asynchronous events, such as
  - Occurrence of an event (timer expires, alarm, etc.,)
  - a user quota exceeds (file size, no of processes etc.,)
  - an I/O device is ready
  - encountering an illegal instruction
  - a terminal interrupt like Ctrl-C or Ctrl-Z.
  - some other process send ( kill -9 pid)
- Each signal has a unique number and represented by macros, which are defined in **<signal.h>**.

# Signals

Signal	Value	Action*	Comment
<b>SIGHUP</b>	<b>1</b>	<b>Term</b>	<b>Hang up</b>
<b>SIGINT</b>	<b>2</b>	<b>Term</b>	<b>Interrupt from keyboard (Ctrl + C)</b>
<b>SIGQUIT</b>	<b>3</b>	<b>Core</b>	<b>Quit from keyboard (Ctrl + \)</b>
<b>SIGILL</b>	<b>4</b>	<b>Core</b>	<b>Illegal Instruction</b>
<b>SIGABRT</b>	<b>6</b>	<b>Core</b>	<b>Abort signal from abort(3)</b>
<b>SIGKILL</b> *	<b>9</b>	<b>Term</b>	<b>Kill signal</b>
<b>SIGSEGV</b>	<b>11</b>	<b>Core</b>	<b>Invalid memory reference</b>
<b>SIGPIPE</b>	<b>13</b>	<b>Term</b>	<b>Broken pipe: write to pipe with no readers</b>
<b>SIGALRM</b>	<b>14</b>	<b>Term</b>	<b>Timer signal from alarm(2)</b>
<b>SIGTERM</b>	<b>15</b>	<b>Term</b>	<b>Termination signal</b>
<b>SIGCHLD</b>	<b>17</b>	<b>Ignore</b>	<b>Child stopped or terminated</b>
<b>SIGSTOP</b> *	<b>19</b>	<b>Term</b>	<b>Stop process</b>

\* Comments in the next slide

# Signals (Contd.).

---

- The default action for a signal can be one of the following, depending on the signal (refer to the previous slide)
  - Term    □ application terminates
  - Core    □ application terminates with core dump
  - Ignore □ application ignores the signal by default
- Note that **SIGSTOP** & **SIGKILL** can not be caught or ignored

# Signals (Contd.).

---

- A process can receive signals from different sources
  - **from operating system**
    - Due to hardware exceptions which trigger signals (e.g., division with zero, invalid memory reference)
  - **from another process**
    - When the other process makes kill() system call
  - **from itself**
    - alarm(), which triggers SIGALRM
- The difference between an exception and a signal is that exceptions are synchronous whereas signals are asynchronous.
  - Some exceptions trigger signals. (e.g., floating point exception generates SIGFPE)



# Sending Signal to a Process

- One process can send signal to another process using **kill** command.

```
$ kill -s signumber processid
```

or

```
$ kill -sname processid
```

```
$ sleep 1000&  
[1] 17719  
$ kill -s SIGKILL 17719
```

```
$ sleep 1000&  
[1] 17814  
$ kill -s 9 17814
```

```
$ kill -TERM 2340 □ request for termination of process with PID  
2340
```

# Sending Signals from program

- One process can send signal to another using **kill()** system call.

```
#include <signal.h>
int kill(pid_t pid, int signum);
```

sends **signum** signal to the **pid** specified

- One process can send signal to itself using **raise()**

```
#include <signal.h>
int raise(int signum);
```

sends **signum** signal to itself

# signal function

```
#include <signal.h>
```

```
typedef void (*sighandler_t)(int);
```

```
sighandler_t signal(int signum, sighandler_t handler);
```

- Handler argument can be one of the following
  - **programmer defined handler** □ to perform specific action
  - **SIG\_IGN** □ to ignore the signal
  - **SIG\_DFL** □ to perform default action (in most cases terminate process)
- returns previous value of the signal handler
- This can be saved, to restore to the old value if signal handler is to be restored

# signal handler

```
#include <signal.h>
void sighandler(int sig)
{
    printf("received signal = %d\n", sig);
}

int main (void)
{
    signal (SIGINT, sighandler);
    while(1)
    {
    }
}
```

Comment the **signal()** call, and run the program.  
Try Ctrl+C and Ctrl+\ while the program is in execution.

Uncomment the **signal()** call, and run the program.  
Try Ctrl+C and Ctrl+\ while the program is in execution.

What is the difference noticed?

# Inheriting signal handlers by child processes

```
#include <signal.h>
void sighandler(int sig)
{
    printf("pid %d received signal
           = %d\n", getpid(), sig);
}
int main (void)
{
    printf("parent pid = %d\n",
           getpid());
    signal (SIGINT, sighandler);
    if (fork() == 0) {
        printf("child pid =%d\n",
               getpid());
    }
    while(1)
    {
    }
}
```

**Does a child process  
inherit signal handler?**

**Run the test program.**

# Handling SIGCHLD to Avoid Zombies

```
void sighandler(int sig)
{
    printf("pid %d received signal = %d\n", getpid(), sig);
}
int main (void)
{
    printf("pid = %d\n", getpid());
    signal (SIGCHLD, sighandler);
    if (fork() == 0) { printf("child pid =%d\n", getpid()); }
    while(1)
    {
    }
}
```

## NOTE:

Just catching of **SIGCHLD** does not ensure that all zombie processes are handled by the parent process.

```
$ ./a.out&
32051
pid = 32051
SIGCHLD = 17
$ child pid =32052
$
$ kill -9 32052
$ pid 32051  received signal = 17
$
$ ps -S
  PID TTY          STAT TIME COMMAND
 32051 pts/1        R    0:20 ./a.out
 32052 pts/1        Z    0:18 [a.out <defunct>]
$
```

# Handling SIGCHLD to Avoid Zombies (Contd.).

```
#include <signal.h>
void sighandler(int sig)
{
    int wstatus;
    printf("pid %d received signal
           = %d\n", getpid(), sig);
    wait(&wstatus);
}

int main (void)
{
    printf("pid = %d\n", getpid());
    signal (SIGCHLD, sighandler);
    fork();
    while(1)
    {
    }
}
```

```
$ ./a.out&
32803
pid = 32803
$
$ ps
32803 pts/1      00:00:06 a.out
32804 pts/1      00:00:06 a.out
$
$ kill -9 32804
$ pid 32803  received signal = 17
$
$ ps
32803 pts/1      00:00:06 a.out
$
```

# alarm () and pause()

- **alarm()** is used to set an alarm for delivering SIGALARM signal.

**unsigned int alarm (unsigned int seconds);**

- On success it returns the number of seconds remaining until previously set alarm due; zero if no alarm is scheduled.
  - **alarm()** is not a blocking call.
  - To cancel an existing alarm, pass value zero (0) as argument.
- **pause()** is a C Library function, which suspends calling process till a signal is caught and the signal handler is returned.

**int pause (void);**

the function returns -1 and sets **errno** to **EINTR**



# alarm () and pause() – example

```
int main (void)
{
    printf("setting alarm(5)\n");
    alarm(5);
    printf("alarm set\n");
    printf("process paused\n");
    pause();
    return 0;
}
```

Note that signal handler is not defined.

Program output

```
$ ./a.out
setting alarm(5)
alarm set
process paused
Alarm clock
$
```

# Alarm with Signal Handler

```
#include <signal.h>
#define ALARM_TIME    5
int sec = 0;
void sighandler(int sig)
{
    printf("received signal =
           %d\n", sig);
    sec += ALARM_TIME;
    printf("%d seconds elapsed\n",
           sec);
    alarm(ALARM_TIME);
}
```

```
int main (void)
{
    signal (SIGALRM,
            sighandler);
    printf("setting alarm\n");
    alarm(ALARM_TIME);
    printf("alarm set\n");
    while(1)
    {
        pause();
    }
}
```

# Alarm with Signal Handler (Contd.).

Program output with signal handler defined

```
$ ./a.out
setting alarm
alarm set
received signal = 14
5 seconds elapsed
received signal = 14
10 seconds elapsed
received signal = 14
15 seconds elapsed
received signal = 14
20 seconds elapsed
...
```

Output continues infinitely as shown above

# POSIX Style Signal Handling

---

- POSIX signal handling has some advantages
  - Reliable signal handling
  - Large number of signals can be handled
  - Allows signals to be blocked and unblocked
  - Possible to define signal sets
- Signal handler is defined using **sigaction()**
- Why to prefer **sigaction()**
  - depending on the variant of Unix, semantics of **signal()** can be C library function or System V or BSD, resulting in portability issues.

# sigaction

---


```
#include <signal.h>
int sigaction( int signum,
               const struct sigaction *act,
               struct sigaction *oldact);
```

- **signum**, is a specified signal
- **act** is used to set the new action of the signal signum;
- **oldact** is used to store the previous action, usually NULL.

# Handling Signal – POSIX way

```
#include <signal.h>
void sighandler(int sig)
{
    printf("received signal = %d\n", sig);
}

int main (void)
{
    struct sigaction action = { sighandler };
    sigaction(SIGINT, &action, NULL);
    while(1)
    {
    }
}
```



# Resource Limits



# Resource Limits

---

- The OS imposes limits for certain system resources it can use.
- Applicable to a specific process.
- The “**ulimit**” shell built-in can be used to set/query the status.
- “**ulimit -a**” returns the user limit values



# Resource Limits (Contd.).

---

- The command **ulimit** takes the following options
  - f     □ Maximum size of the files created.
  - l     □ Maximum amount of memory that can be locked using **mlock()** system call.
  - n     □ Maximum number of open file descriptors.
  - s     □ Maximum stack size allowed per process.
  - u     □ Maximum number of processes available to a single user.

# Resource Limits (Contd.).

---

- Resource limits are of two types
  - Hard limit
    - Absolute limit for a particular resource. It can be a fixed value or “unlimited”
    - Only super user can set hard limit.
    - Hard limit once set, cannot be increased.
  - Soft limit
    - User-definable parameter for a particular resource.
    - Can have a value from 0 till <hard limit> value.
    - Any user can set soft limit.
- Limits are inherited (the new values are applicable to the descendent processes).

# getrlimit/setrlimit

- `getrlimit()/setrlimit()` lets an application get/set resource limits

```
#include <sys/time.h>
#include <sys/resource.h>
int getrlimit(int res, struct rlimit *reslimit);
int setrlimit(int res, const struct rlimit * reslimit);

struct rlimit {
    rlim_t rlim_cur;    /* Soft limit */
    rlim_t rlim_max;    /* Hard limit */
};
```

Note: soft limit value must be less than hard limit value

# getrlimit/setrlimit (Contd.).

---

Resource that can be set with limits

- `RLIMIT_FSIZE`      □ Maximum size of the file.
- `RLIMIT_MEMLOCK` □ Maximum amount of memory that can be locked.
- `RLIMIT_NOFILE`    □ Maximum number of open file descriptors.
- `RLIMIT_STACK`     □ Maximum stack size allowed.
- `RLIMIT_NPROC`     □ Maximum number of process available to a single user.

# setrlimit() - example


```
#include <fcntl.h>
#include <sys/resource.h>
main(int argc, char *argv[])
{
    int fd, k;
    char *buf = "hello ";
    struct rlimit rlim = {30,100 };
    setrlimit(RLIMIT_FSIZE, &rlim);
    fd = open(argv[1], O_CREAT | O_TRUNC | O_WRONLY, 0600);
    for (k=0; k<20; k++)    {
        printf("attempting to write  %d bytes\n",
strlen(buf));
        write(fd, buf, strlen(buf));
        printf("successfully written %d bytes\n",
strlen(buf));
    }
    close(fd);
}
```

# setrlimit() – example (Contd.).

---

Run the previous program and observe that the program terminates after writing 30 bytes.

```
$ ./a.out out
attempting to write 6 bytes
successfully written 6 bytes
attempting to write 6 bytes
successfully written 6 bytes
attempting to write 6 bytes
successfully written 6 bytes
attempting to write 6 bytes
successfully written 6 bytes
attempting to write 6 bytes
successfully written 6 bytes
attempting to write 6 bytes
File size limit exceeded
$
```



# Hands-on and Assignments



# Exercises

---

- 1) Define a signal handler that catches SIGINT, SIGTERM and SIGQUIT, prints the signal it has received.  
Test the program that it functions as desired.
  
- 2) Write a program which meets the following requirements.
  - a) Parent process registers a signal handler for SIGINT, which prints process ID and the signal which the handler received.
  - b) Parent process creates two child processes both run in infinite loop.
  - c) Parent process sends SIGINT to both child processes
  - d) Parent process sends SIGINT to itself

Run the program and observe the behaviour.

Note: Use POSIX sigaction to register signal handler.



# Exercises (Contd.).

---

- 3) Write a program to check if resource limits set by a process are inherited by its child process.
- 4) Write a program to limit the number of files that can be opened by a process to not more than 10 files and test if that works. Can this be increased subsequently to 15 and open 5 more files?

# Summary

---

In this module, we discussed:

- the purpose of signals
- how to program signals
- POSIX signal handling
- various resource limits in UNIX
- how to alter resource limits
- impact of resource limits on processes

# Review Questions


---

- a) Is signal handling synchronous or asynchronous?
- b) For which signals, signal handler can not be defined?
- c) What are the advantages of POSIX signal handling?
- d) What should be done to ensure that a terminated process does not end up as a zombie?
- e) For each of the following identify whether it is a blocking call or not.
  - i) pause
  - ii) alarm
- f) Does a process inherit signal handlers defined by its parent?
- g) What is the significance of resource limits?
- h) Does a process get terminated if it exceeds resource limit?

# References

---

- 1) W. Richard Stevens and Stephen A. Rago. Advanced Programming in the UNIX Environment. Ed 2, US: Addison-Wesley, 2009.
- 2) Kay A. Robbins and Steven Robbins, UNIX Systems Programming, New Delhi: Pearson Education, 2009.
- 3) Rochkind, Advanced Unix Programming, Ed 2. US: Addison-Wesley, 2004.
- 4) Arnold Robbins, Linux Programming by Example, New Delhi: Prentice Hall, 2008.



**Thank You**